Proofed: concept note

One small ASR model that writes clean, proofed text: integrated ASR and post-processing

Daniel Rosehill · 24 October 2026

Project	github.com/danielrosehill/Proofed (public)
Status	Phase 1 research done (viable). No training run yet.
Licence intent	Non-commercial research prototype, published openly on Hugging Face
Source of truth	The 24-minute project voice note (context/voice-notes/2026-10-24-proofed-brief/), then docs/development-plan.md

Summary

The idea: fine-tune a stock speech-recognition model (NVIDIA Parakeet TDT 0.6B v2 first) on post-processing tasks only, so it writes clean, proofed text straight from audio. No ums, repeats or false starts; numbers as digits; self-corrections applied; paragraph breaks; remarks to someone else in the room left out; whispered speech handled.

The headline: transcription stays the core task. Proofed is not a large audio-multimodal model that can do everything and is heavy because of it. It is a small, specialised ASR model that also does the post-processing every user wants, in the same model and ideally the same pass, instead of a second model afterwards.

What we would value your view on: whether the transducer can learn these edits well (§5), whether one pass or two passes inside one checkpoint is the right architecture (§5), and how best to add paragraph breaks to Parakeet (§6). Questions are collected in §8.

1 The idea

Voice dictation today is a two-model job. An ASR model writes what was said, word for word, and a second model (a regex layer, a small text model or an LLM) cleans it up: it removes fillers, applies “scratch that”, writes numbers as digits and adds paragraphs. The cleanup that matters to nearly everyone is the same every time, yet it costs a second model, a second inference step and, on a phone, a second download.

Proofed tests a simple question: **how much of that general-purpose cleanup can be baked into the ASR model itself?** The base is an unmodified Parakeet TDT 0.6B v2. The fine-tune is trained only on audio paired with the clean target, task by task, each task evaluated on its own. Domain vocabulary (names, technical terms, code-switching) is deliberately excluded: it is a separate layer that can be fine-tuned on top later.

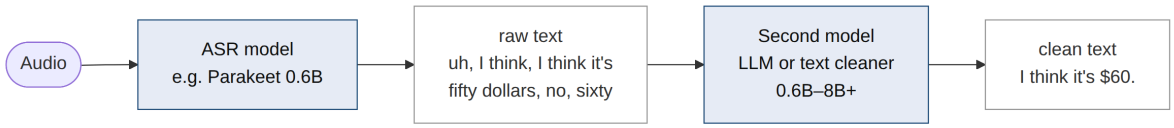


Figure 1: Today: two models in series. Proofed folds the second into the first.

1.1 Why integrate, and why not just use an audio LLM

	ASR + LLM cleanup (today)	**Audio-multimodal LLM**	**Proofed**
Core task	transcription, then text editing	everything (chat, reasoning, audio QA)	**transcription, written cleanly**
Models at inference	2	1, large	**1, ~0.6B**
Size on device	ASR + 0.3–8 GB text model	several GB and up	**ASR only**
Latency	ASR plus a full LLM pass	slow autoregressive decode	**ASR decode, plus at most a small second head**
Hallucination risk	LLM can rewrite content	high, open-ended output	**low: output is constrained to what was said**
Instruction-following	prompt-driven, variable	prompt-driven, variable	**fixed: the same cleanup every time**
Sees the audio when cleaning	no (text only)	yes	**yes**

The point of the last row: a text-only cleaner cannot hear a pause, a change of voice or a whisper. A cleanup step inside the ASR model can use prosody to decide where a paragraph ends or that a remark was aimed at someone else.

2 Challenges identified

From the voice note and the phase 1 research (report in the repo, 125 sources):

1. **Nobody has done this on a transducer.** Disfluency removal end-to-end is published for CTC, attention and RNN-T models on Switchboard, and with a pretrained acoustic model it matched or beat ASR + cleanup (Bassi et al. 2023: 12.2% vs 13.1% WER). There is no Parakeet or TDT version, and nothing on dictation self-corrections, same-speaker asides or paragraphs.
2. **Some edits look ahead or reorder.** Deleting “Haifa” in “Haifa today, I mean Tiberias” is only decidable after hearing “I mean”. “fifty dollars” → “$50” emits a symbol before the word it came from. A monotonic transducer may find these hard (§5).
3. **Paragraph breaks need a token Parakeet does not have.** The v2 tokenizer is a 1,024-piece SentencePiece BPE with `nmt_nfkc` normalisation, which turns `\n` into a space (§6).
4. **Catastrophic forgetting.** A narrow full fine-tune has been reported to move WER from 5.12% to 51.08%; replay data or joint adapters held it at 5.79% (research §3.4).

5. **Data.** About 42 minutes of in-scope recordings exist today, none yet for self-corrections. Acted disfluency (read from a prompt) may not transfer to natural dictation: untested anywhere.
6. **Where the line is.** Only transformations that are right for every user every time belong in the model. Anything use-specific (email formatting, style) stays outside.

3 The proposal

Fine-tune Parakeet TDT 0.6B v2 on (audio, clean text) pairs drawn from a fixed taxonomy of post-processing tasks, with replay data to protect plain transcription, and score every task separately: task success, WER on the words that should be kept, and regression on plain English (LibriSpeech-PC). Start with the core tier, then add the ambitious and experimental tiers as separate measured experiments, so that a task that cannot be learned at this scale is identified rather than hidden in an average.

Data: own read-aloud recordings per task (the recorder app has a dedicated Proofed mode; target about 4 hours for the first run), redistributable public corpora (DisfluencySpeech, AMI, ICSI, LibriSpeech-PC for replay), and TTS synthesis for self-corrections, asides and the long tail of numbers.

4 The integrated post-processing

Each task is classed by what it does to the token sequence, because that decides which decoder can learn it (§5). Full version: `docs/taxonomy.md`.

Id	Task	Op	Said → written	Tier
A1	Filled pauses	delete	“I, uh, think so” → “I think so”	core
A2	Repetitions	delete	“the the report” → “the report”	core
A3	False starts	delete	“And going to the — I’m going to the cinema” → “I’m going to the cinema”	core
A4	Abandoned sentences	delete-ctx	“I think I had an idea — I want to go to the movies” → “I want to go to the movies”	core
B1	Self-corrections	delete-ctx	“Haifa today. I mean Tiberias” → “Tiberias today”	ambitious
B2	Editing commands	delete-ctx	“Jerusalem. Scratch that, Tel Aviv” → “Tel Aviv”	ambitious
C1	Numbers, dates, units	rewrite	“twenty twenty-six” → “2026”	core
C2	Money	reorder	“fifty dollars” → “$50”	core
D1	Paragraph breaks	structure	(pause, topic shift) → blank line	ambitious
E1	Asides to someone else	filter	“the docs are here — just leave it by the door — now” → “the docs are here now”	experimental
F1	Whispered speech	robust	(whispered) → same text	robustness

Operation types: *delete* drops a span in place; *delete-ctx* drops a span whose deletion depends on later audio; *rewrite* changes the written form in order; *reorder* moves a token; *structure* inserts layout; *filter* drops speech not meant for the transcript; *robust* changes only the acoustic condition. Punctuation and casing are inherited from the base model and must survive.

Negative controls sit inside each task's prompts: intentional repetition ("very, very good") stays, and numbers that are names or idioms ("Formula One") stay words. Over-deletion is the failure the evaluation watches.

5 Architecture: one forward pass or two

The voice note asks whether one model can "encapsulate two passes, the first pure ASR, the second post-processing". Four options, from least to most change to Parakeet:

5.1 Option A: single pass, fine-tuned TDT

Keep the architecture. Fine-tune the prediction network and joint (and optionally the upper encoder layers, or adapters) on clean targets. Deletion becomes "emit blank over the span"; rewriting is monotonic. The FastConformer encoder is offline and full-context, so the encoder frame for "Haifa" already carries information about the "I mean" that follows.



Figure 2: Option A: single pass. Same model, new targets.

For: zero inference cost, zero architecture change, the existing NeMo recipe and every existing export (ONNX, sherpa-onnx) still work. **Against:** the decoder still commits left to right, so delete-ctx and reorder tasks depend on the encoder having learned to encode the decision early; TDT duration prediction over long deleted spans is untested.

5.2 Option B: two passes inside one checkpoint (deliberation)

Pass 1 is the existing TDT head, producing the verbatim transcript. Pass 2 is a small Transformer decoder that attends to both the encoder states and the pass-1 tokens, and writes the clean text. Both share one encoder, ship as one model file and run in one call. This is the Google two-pass and deliberation line of work: LAS rescoring of RNN-T (Sainath et al. 2019, arXiv 1908.10992; 17–22% WER reduction, under 200 ms added latency) and deliberation (Hu et al., ICASSP 2020, arXiv 2003.07962; a further 12% relative over LAS rescoring).

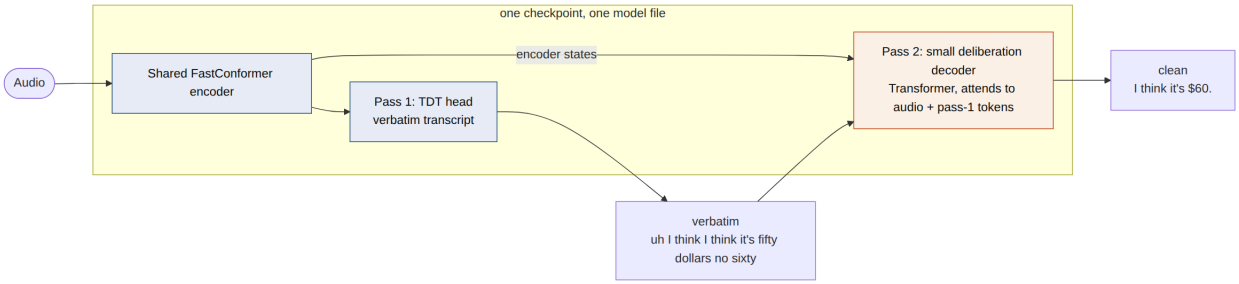


Figure 3: Option B: two passes, one checkpoint. The second pass sees the audio and the whole verbatim draft.

For: pass 2 sees the whole utterance, so look-ahead edits and reordering become ordinary sequence-to-sequence work; pass 1 keeps an exact verbatim output for free (useful for evaluation and for users who want it); the deliberation decoder can be small (tens of millions of parameters). **Against:** new

module and training code in NeMo, a second autoregressive decode, and export work for on-device use.

5.3 Option C: task-token attention decoder (Canary-style)

An attention encoder-decoder with a task token (`<verbatim>` or `<clean>`) prompting the decoder, as Canary does for transcription and translation. One pass, any edit, but it changes the base model family (a candidate is canary-180m-flash), and attention decoders hallucinate more than transducers on long or noisy audio.

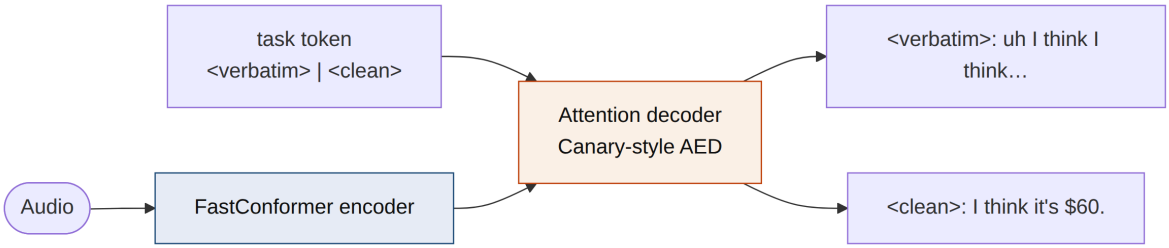


Figure 4: Option C: one AED, two output modes selected by a task token.

5.4 Option D: tag, then delete

The decoder emits verbatim text with tags around spans to drop (Futami et al.; Mendelev et al. found tagging better than direct deletion), and a trivial rule removes them. It keeps a verbatim record and makes deletions inspectable, but it does not help rewriting or reordering.

5.5 Comparison and recommended path

	A: single pass	**B: two-pass, one checkpoint**	**C: task-token AED**	**D: tag-then-delete**
Architecture change	none	add a small decoder	different base	none (targets only)
Inference cost	1×	1× encoder + small decoder	1× AED decode	1×
delete	good	good	good	good
delete-ctx	uncertain	good	good	uncertain
reorder ($50)	uncertain	good	good	no help
Verbatim kept	no	yes	yes (other mode)	yes
On-device export	existing	new work	existing (Canary)	existing

Plan: run A first, task by task, because it costs nothing at inference and answers the core question directly. Add B only for the task families where A fails (expected: B1/B2 and C2). Run C with canary-180m-flash as a control, so a failure of A can be attributed to the decoder type rather than to the data.

6 Paragraph breaks

Paragraphs matter to the project more than any other formatting task, and no ASR model was found that emits them natively; AssemblyAI and Deepgram add them after recognition, on the API side.

What Parakeet TDT 0.6B v2 offers: no newline piece. A sweep of all 22 NVIDIA Parakeet, Canary and Nemotron speech tokenizers (pulled from the `.nemo` archives, with `\n`, `\n\n`, ¶ and U+2029 round-tripped) found that every one strips the newline. The only partial exception is canary-qwen-2.5b, whose Qwen3 decoder has newline tokens, at 2.5B parameters.

Route	How	For	Against
1. Add a paragraph token to v2 (recommended)	Add `"\n"` (or a named para token) to the tokenizer, as NVIDIA did for `<EOU>`/`<EOB>` in parakeet_realtime_eou_120m-v1 (ids 1024–1025); warm-start the joint so the new row sits at 1024, blank moves to 1025, and the five TDT duration rows stay last	a real token; `\n\n` round-trips exactly (tested); all pretrained weights kept	NeMo's warm-start routine was written for plain RNN-T and needs adapting for TDT's row layout (read from code, not yet run)
2. ^ sentinel (control arm)	Train with ^ (id 918, one token) at paragraph boundaries; replace after decoding	no model change at all	borrows a real character; a hack
3. v3 reserved token (fallback)	parakeet-tdt-0.6b-v3 ships 30 unused reserved tokens (spltoken0–29) the model can already emit	no code change	leaves v2, which is better in English

Avoid NeMo `change_vocabulary`: it reinitialises the decoder and joint and drops the TDT loss. Whether the TDT duration head learns to emit a paragraph token cleanly at pauses is untested. Full survey: `research/2026-10-24-phase-1-feasibility/notes/06-parakeet-newline-survey.md`.

7 Data and plan

Phase	What	Status
1	Research: prior art, paragraph token, base models, feasibility, datasets	done 2026-10-24
2	Data: record per task (≈4 h target), public corpora, TTS synthesis, publish the blend	in progress
3	Evaluation: per-task test sets, natural-dictation eval (40 spontaneous dictations)	planned
4	Prototype: Option A fine-tune on the core tier, then the other tiers	planned
5	Comparisons: Option B where A fails; Canary and Qwen3-ASR-0.6B arms	planned

The training data so far is published as a public dataset, `danielrosehill/Proofed-ASR-Cleanup-Training-Data` (CC-BY-4.0).

8 Questions for the Parakeet team

1. **TDT and deletion.** Have you seen TDT models learn to drop long spans (fillers, false starts) cleanly? Does the duration head cope, or does it push towards skipping frames in ways that hurt the kept words?
2. **Look-ahead edits.** For self-corrections ("Haifa, I mean Tiberias"), would you expect the offline encoder to carry enough right context for a single-pass transducer, or is a second pass (Option B) the realistic route?
3. **A paragraph token in v2.** Is adding a token the way parakeet_realtime_eou_120m-v1 added `<EOU>`/`<EOB>` the right route for a TDT model, and is our row layout (new token at 1024, blank after it, durations last) correct? Is a newline or paragraph token planned for future models?
4. **Fine-tuning recipe.** For narrow tasks with a few hours of data: full fine-tune with replay, or adapters in the joint and decoder? Which layers would you unfreeze?
5. **Deliberation in NeMo.** Is there an existing NeMo component (or internal experience) for a second-pass decoder on top of a transducer that you would build on?
6. **Anything we have missed:** prior work at NVIDIA on clean or formatted output from Parakeet or Canary.

9 Links

- Repository: https://github.com/danielrosehill/Proofed
- Project voice note (MP3, 24:40): https://github.com/danielrosehill/Proofed/blob/main/context/voice-notes/2026-10-24-proofed-brief/voice-note.mp3
- Phase 1 research report (PDF): https://github.com/danielrosehill/Proofed/blob/main/research/2026-10-24-phase-1-feasibility/Proofed-research-report.pdf
- Task taxonomy: https://github.com/danielrosehill/Proofed/blob/main/docs/taxonomy.md

Proofed is a non-commercial research project. Models and data will be published openly.